

Hardware-Assisted Memory Safety for WebAssembly

Fritz Rehde

Advisor: Martin Fink

Chair of Distributed Systems and Operating Systems

<https://dse.in.tum.de/>



15.04.23 – 15.08.23

- “Binary instruction format for a stack-based virtual machine”¹
- “Compile once, run everywhere” (**Portable**)
- Environments: (**Portable**)
 - Web browser (V8)
 - Systems programming (Wasmtime) with WASI
- Compilation target for low-level programming languages, e.g. C (**Performant**)
- Sandboxing: protect host though bounds checking inside linear memory (**Secure**)
 - Does not protect module from itself within sandbox!


¹ <https://webassembly.org/>

- Program is memory safe if all memory pointers refer to valid memory when being dereferenced

Spatial: buffer overflows



Temporal: use-after-free, double-free



How can we provide memory safety in WebAssembly without incurring significant performance costs?

System design goals:


- **Safety:** Memory safety issues should cause runtime crashes instead of UB
- **Performance:** Minimal performance overhead
- **Usability:** No need to modify input C source code

ARM64 Hardware Extensions

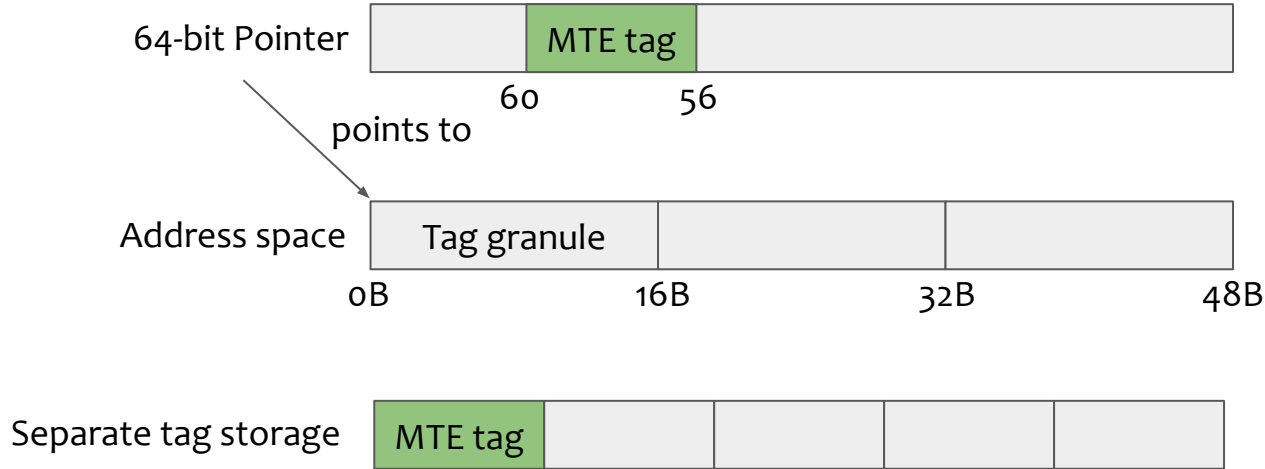
- Store metadata in unused upper bits in pointers

- ~~Motivation~~
- Background
 - Memory Tagging Extension (MTE)
 - Pointer Authentication (PAC)
- Design
- Evaluation

- Memory Tagging Extension (MTE) in ARMv8.5:
 - Ensures spatial and temporal memory safety
 - Pointers and memory locations are tagged (“colored”)
 - Tags are compared on memory access




ARM64 Hardware Extensions: MTE



<code>IRG(ptr) -> tagged_ptr</code>	"Insert Random Tag" into pointer
<code>STG(tagged_ptr, ptr)</code>	"Store Allocation Tag" in memory at 16-byte granularity

- Pointer Authentication in ARMv8.3:
 - Protects against malicious pointer overrides (e.g. ROP, JOP)
 - Pointer Authentication Code (PAC): cryptographic signature embedded in pointer



with $y - x = 55 - \text{<linux virtual address size>}$ ¹

<code>PACDA(ptr) -> signed_ptr</code>	Generate PAC, and sign pointer with it
<code>AUTDA(ptr) -> authed_ptr</code>	Authenticate pointer, remove PAC if succeeded

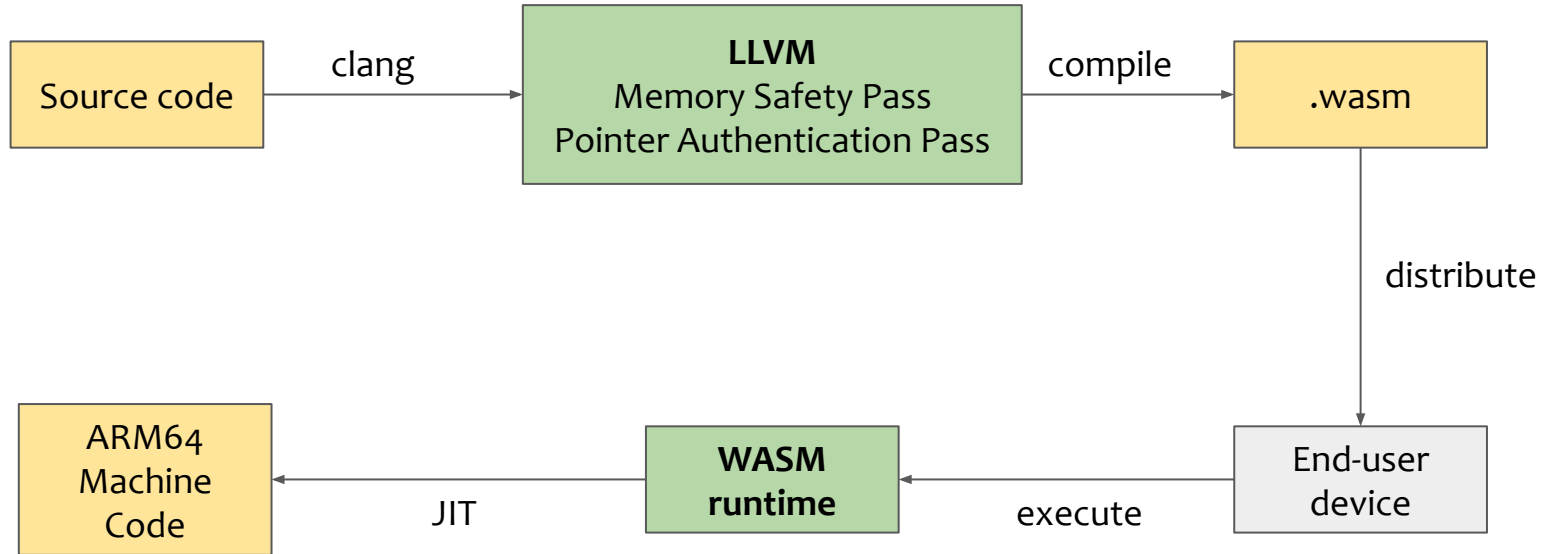
¹ <https://www.kernel.org/doc/html/v6.5/arch/arm64/pointer-authentication.html>

Outline



- ~~Motivation~~
- ~~Background~~
- Design
- Evaluation

System overview



- ~~Motivation~~
- ~~Background~~
- Design
 - New WASM Instructions
 - LLVM Passes
 - Wasmtime Additions
- Evaluation

Design: new WASM instructions

- Introduce *segments*: protected memory region, only accessible with matching pointer

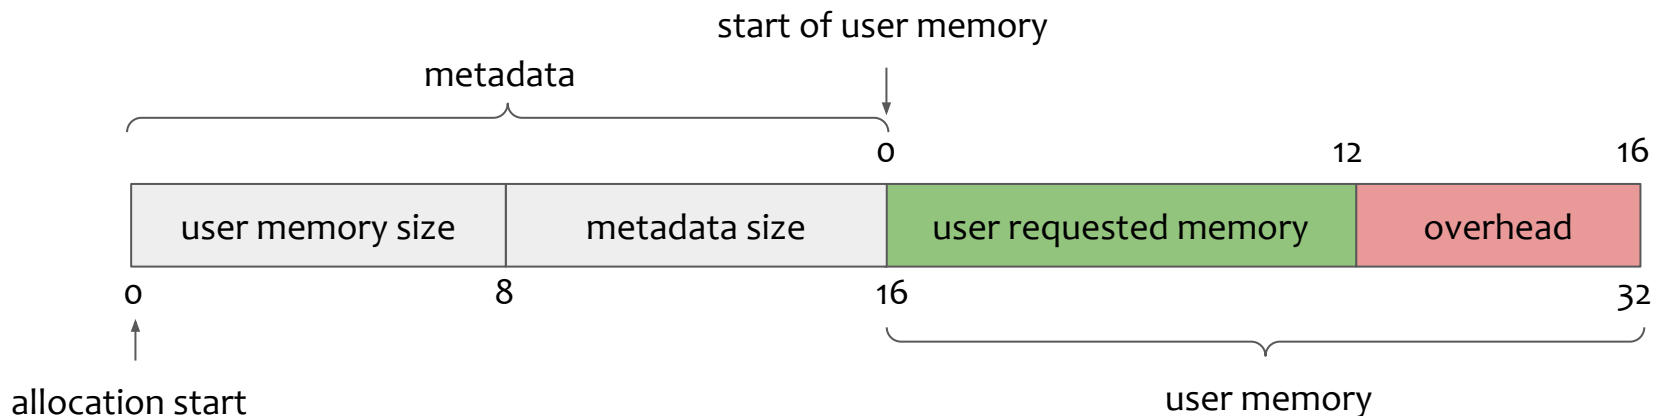
WASM instruction	Implementation in WASM runtime
<code>segment.new(index¹, size) -> tagged_index</code>	<ol style="list-style-type: none">1. Generate new (random) tag (IRG)2. Tag index with new tag3. Tag memory with new tag in 16-byte granules (STG)
<code>segment.free(index, size)</code>	Tag memory with free tag o (STG)
<code>pointer_sign(index) -> signed_index</code>	Sign index (PACDA)
<code>pointer_auth(index) -> authed_index</code>	Authenticate index (AUTDA)

¹ An index is an i64

Design: LLVM Memory Safety pass

- Wrappers around `aligned_alloc`, `malloc`, `calloc`, `realloc` and `free`
- MTE requires 16-byte alignment, but we also need to handle custom alignment

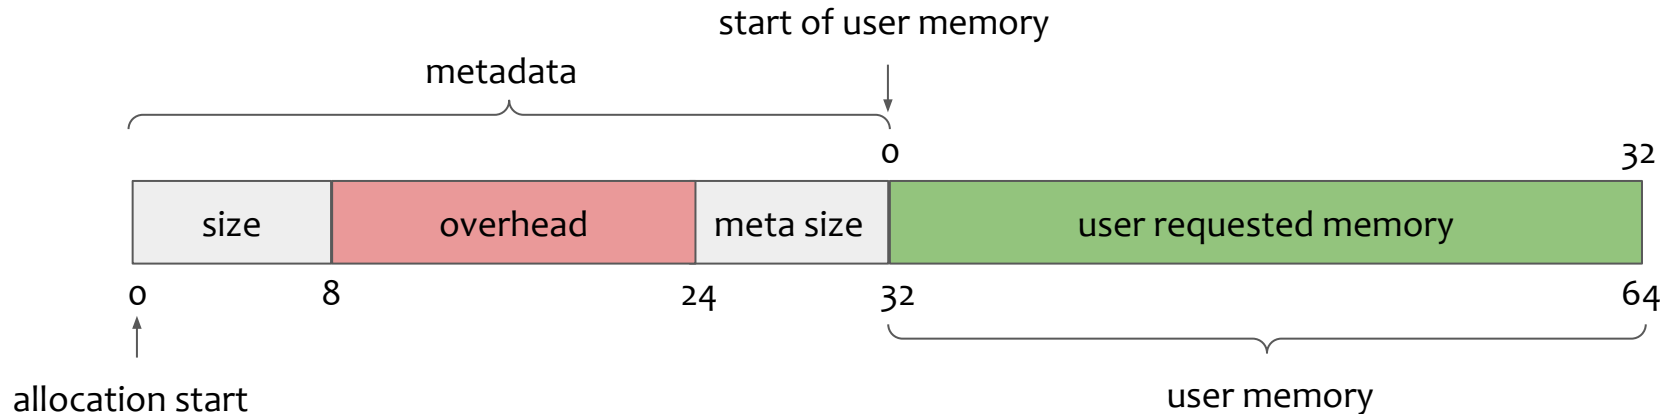
malloc(12)



Design: LLVM Memory Safety pass

- Wrappers around `aligned_alloc`, `malloc`, `calloc`, `realloc` and `free`
- MTE requires 16-byte alignment, but we also need to handle custom alignment

`aligned_alloc(32, 32)`



Design: LLVM Pointer Authentication pass

- Before storing a pointer value, sign it (with `pointer_sign`)
- After loading a pointer value, authenticate it (with `pointer_auth`)
- Rules for implementation:
 - **Signed** pointers **must be** authenticated before being dereferenced
 - **Non-signed** pointers **may not be** authenticated before being dereferenced
- **Problem:**
 - WASI-libc only provides interface (header files)
 - WASM runtime provides implementations
- **Solution:**
 - Cannot sign/authenticate pointers, which come from or are passed to “external functions”
 - External functions: function-, module- and program-granularity (WIP: LTO)

Design: WASM runtime

- Modify WASM runtime *Wasmtime*¹ (Bytecode Alliance; Rust)
- Generate MTE and PAC instructions if target supports them
- **Problem:**
 - Random tag generation: neighbouring segments (stack) might have same tag
 - Buffer overflow not detected (probability: $1/15 = 6.67\%$)
- **Solution**²: (WIP)
 - Use IRG for first segment
 - Generate next tags by incrementing

¹ <https://wasmtime.dev/>

² <https://arxiv.org/abs/2204.03781>

Outline

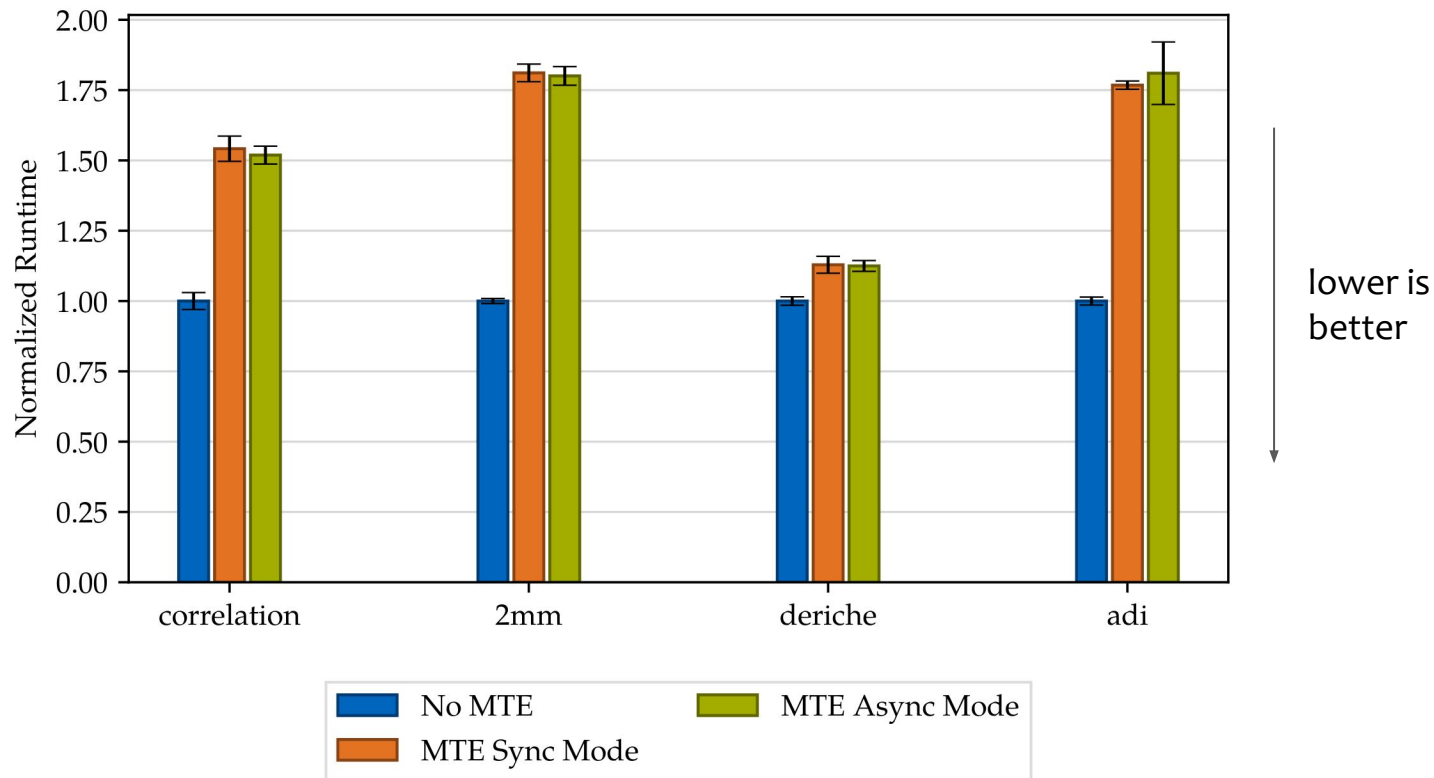


- ~~Motivation~~
- ~~Background~~
- ~~Design~~
- Evaluation

Evaluation: MTE

- Experimental setup:
 - AMD EPYC 7713P CPU (64 cores, 128 threads)
 - 515 GiB DDR4 RAM
 - x86, NixOS 23.05
- MTE not available on real hardware => QEMU


Evaluation: PolybenchC¹



¹ <https://github.com/MatthiasJReisinger/PolyBenchC-4.2.1> (dataset: medium)

Evaluation: Sorting Algorithms

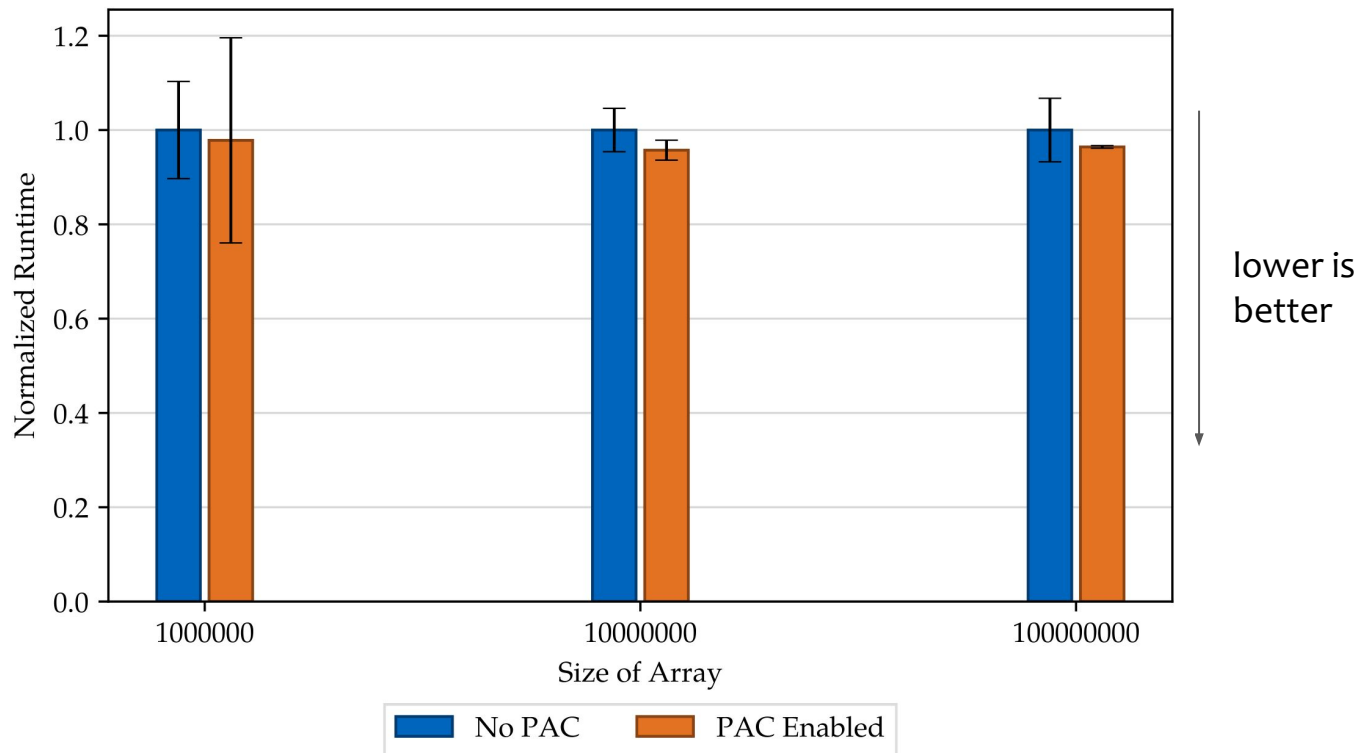
Sort array of size 40,000



- Experimental setup:
 - Apple M1 Pro
 - 10 cores (8 performance; up to 3220 MhZ, 2 low power)
 - 192 KB instruction cache, 128 KB data cache, 24 MB shared L2 cache
 - 16 GB LPDDR5-6400
- Supports PAC natively (through Apple's HVF hypervisor through QEMU)

Evaluation: PAC Overhead

Store n pointers in pointer array (PACDA), load n pointers from pointer array (AUTDA)



- Ensure memory safety in WASM:
 - New memory safety WASM instructions
 - LLVM IR passes to insert them
 - Wasmtime to insert ARM64 MTE and PAC instructions
- PAC: limited safety, but minimal overhead
- MTE: high safety, but real-world performance still unknown

Try it out!

<https://github.com/TUM-DSE/llvm-memsafe-wasm>

<https://github.com/TUM-DSE/wasmtime-mte>

Backup

Example: PAC protection

```
#include <stdio.h>
#include <stdlib.h>

int main() {
    char *string = "Hello World!";
    char **pointer_storage = &string;
    char name[10];

    printf("What is your name?\n");
    scanf("%s", name); // potential buffer overflow
    printf("Hello user %s!\n", name);

    char *loaded_string = *pointer_storage; // failed authentication
    printf("String protected with PAC: %s\n", loaded_string);

    return 0;
}
```


Bottom of the stack (higher memory addresses)

...
string
pointer_storage
name[9]
name[8]
...
name[1]
name[0]
...

Top of the stack (lower memory addresses)

Work in Progress: Optimizing bounds checks using MTE

- Linear memory access: check index is in-bounds, to protect runtime/host
- Our idea to eliminate this overhead using MTE:
 - Tag outside linear memory with **default free tag**
 - Tag entire linear memory and freed user memory with **linear memory free tag**
 - Tag allocations inside with **randomly generated tag**



- Memory region tagging optimizations:
 - ST2G: “Store Allocation Tags”, tag 2 granules at once (32 bytes)
 - Dynamic size: ST2G loop, conditional STG
 - Static size: Loop unrolling threshold (= 160 bytes = 5 x ST2G)
- Error message on MTE trap:
 - Extend signal handler: catch SIGSEGV, check si_code for SEGV_MTESERR or SEGV_MTEAERR

Pointer cannot be signed/authenticated if:

- “Pointer value comes from elsewhere”:
 - Value was passed as parameter to the current function
 - Value was loaded from any memory location
 - Value is a global value
 - Value is return value of any function
- “Pointer value has other uses”:
 - Value is recursively passed as function parameter to an external function
- Aliases have to be checked as well

- Providing Memory Safety in WebAssembly:
 - MSWasm¹: *segments* only accessible through *handles* (fat pointer); software-based; performance overhead
 - GC proposal²: introduce native GC instead of PLs having to ship own GC (large binaries)
- Providing Memory Safety using ARM64 Hardware Extensions
 - HWASan³: software-based tags using Top Byte Ignore (TBI); code size overhead
 - Deterministic Tagging⁴: LLVM analysis to differentiate between safe and unsafe memory allocations => tag only unsafe => performant

¹ <https://arxiv.org/abs/2208.13583>

² <https://github.com/WebAssembly/gc>

³ <https://source.android.com/docs/security/test/hwasan>

⁴ <https://arxiv.org/abs/2204.03781>