

Analysis and Enforcement of GDPR Rules on Key-Value Stores

Ertugrul Aypek

Advisor: Dimitrios Stavrakakis, M.Eng.

Chair of Distributed Systems and Operating Systems

<https://dse.in.tum.de/>

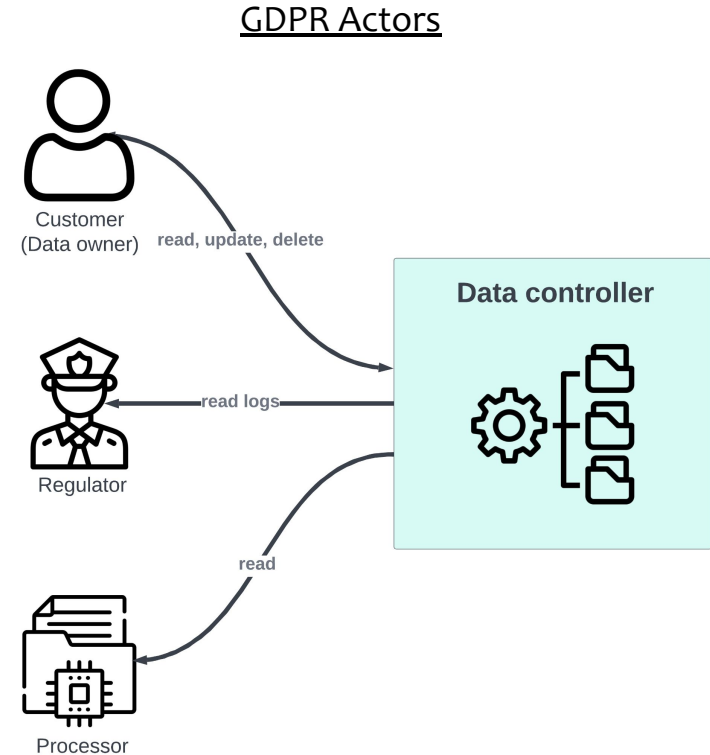


15.02.2023 – 15.08.2023


Background: GDPR definition

General Data Protection Regulation (GDPR)

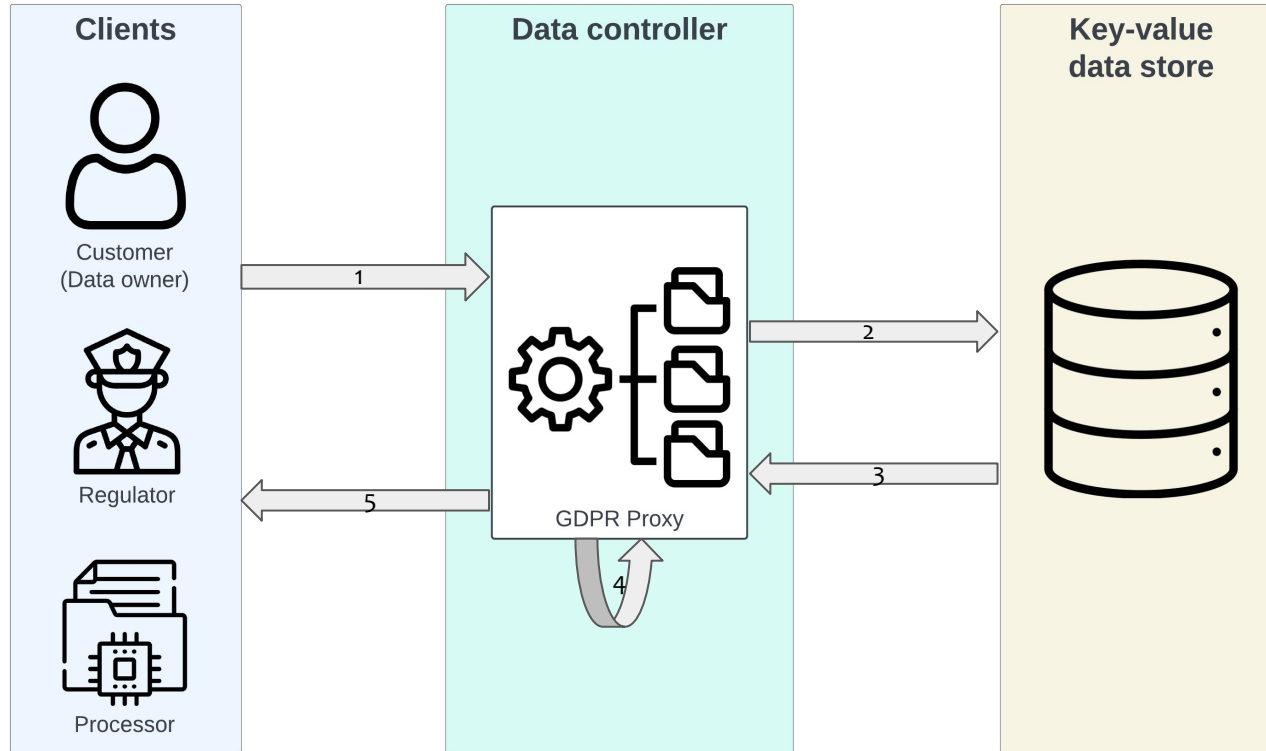
- Introduced in 2018 by EU
- Resulted in more than €2.5B in fines
- Blocks businesses in EU market
- Requires system redesign



Background: GDPR requirements



- For GDPR compliance, each system require internal redesign
- Hard to comply with all the rules and verify compliance
- Performance overhead implications
- Little to no focus on key-value store GDPR compliance in literature



Values are enriched with metadata to be able to comply with GDPR rules:

- user key
- origin
- purpose
- objection
- share
- expiration


Design: Policy Language

- The unified way to interact with the GDPR Proxy
- Default policy as JSON and query overrides
- Policy Compiler to parse queries and enforce query language syntax
- Query Rewriter to merge default policy and query predicates

```
query(PUT("gdpr1","VAL"))&userKey("user1")&origin("src1")  
&objection("purpose3")&purpose("purpose0,purpose1,purpose2")  
      &share("user0")&expiration("o")
```

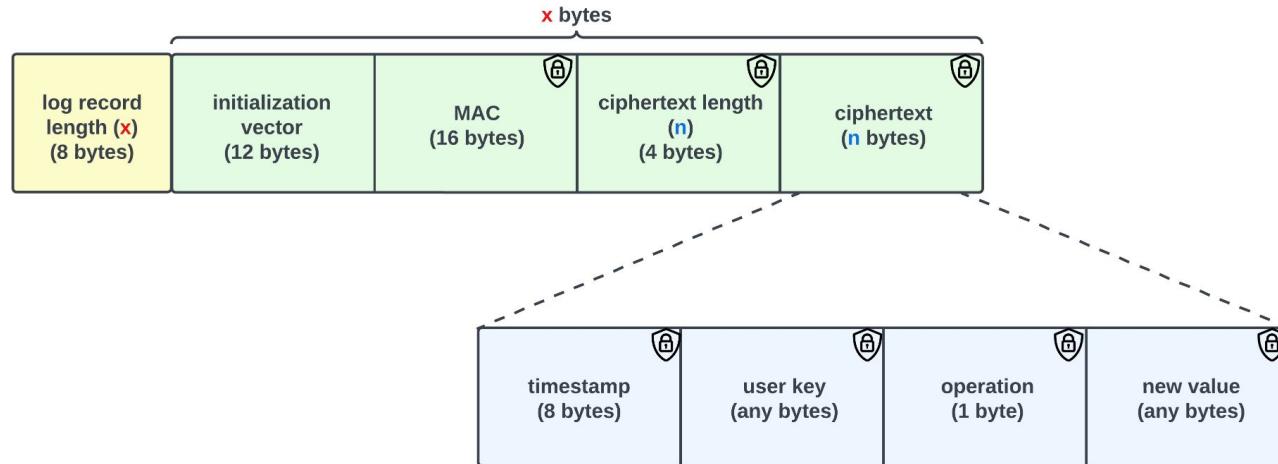
Design: Cipher Engine

- Enables encryption/decryption of metadata enriched values and logs
- Implemented using AES-GCM 128 bit algorithm



Design: Logging Engine

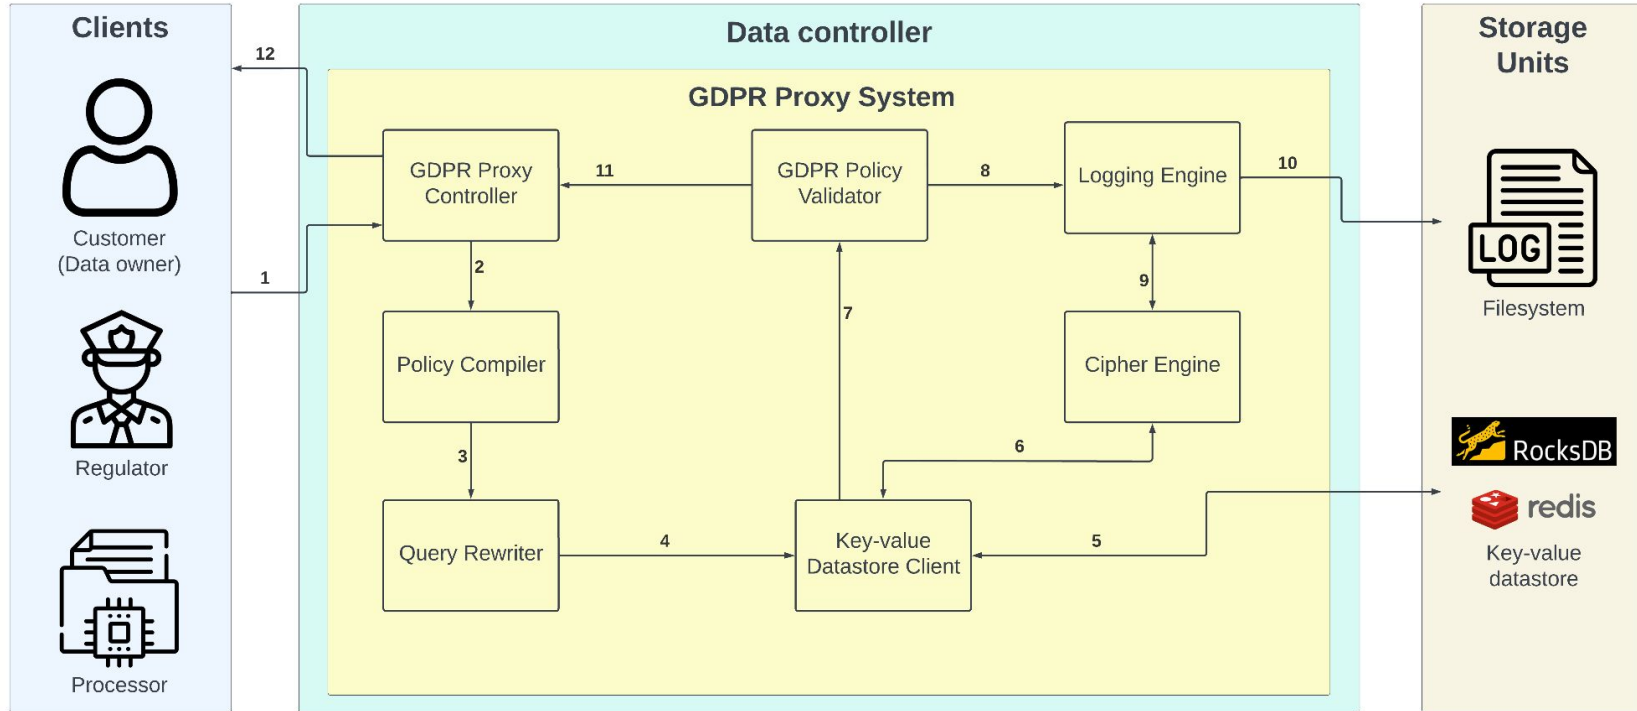
- Used to prove the GDPR compliant processing activities
- Can be kept encrypted in file system



Design: Proxy Controller

- Entrypoint to the system
- Handles user sessions over secure TCP to execute queries
- Parses and stores default policies for each session
- Orchestrates connection to different datastore backends

Design: Design Revisited



Evaluation aspects:

- Correctness (via GDPRBench workloads)
- Speed
- Space

Workloads A-F in GDPRBench with

- 1 million records (via put queries)
- 1 million operations (put/get/delete queries)


Workload	Operation	Application
A	Read/Update (50/50%)	Session store
B	Read/Update (95/5%)	Photo tagging
C	Read (100%)	User profile cache
D	Read/Insert (95/5%)	User status update
F	Read-Modify-Write (100%)	User activity record

Different types of controllers to measure system performance


Functionality	Native Controller	GDPR Controller
GDPR Metadata	no	yes
Policy Language	no	yes
Policy Compiler	no	yes
Query Rewriter	no	yes
Cipher Engine	no	optional
Key-value Client	yes	yes
Logging Engine	no	optional
Policy Validator	no	yes

Evaluation: GDPR Metadata

GDPR metadata overhead with RocksDB



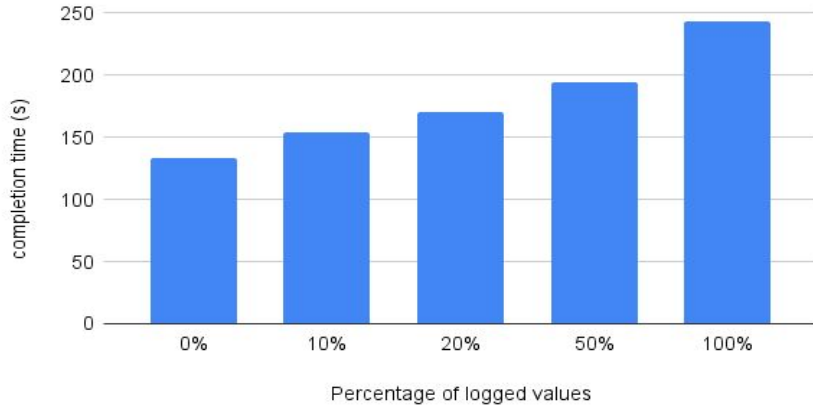
GDPR metadata overhead with Redis



GDPR metadata processing of all workloads is 90% for RocksDB and 73% for Redis


Evaluation: Logging

Logging time overhead with RocksDB



Lower is better

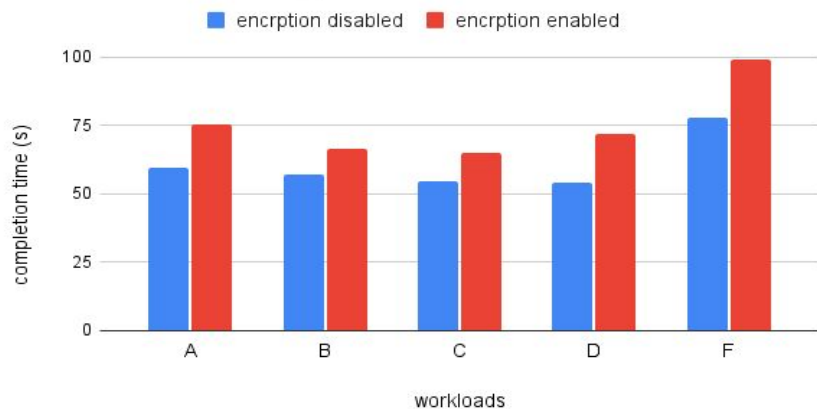
Logging time overhead with Redis



Logging 0% vs 100% of queries adds around 100 seconds overhead

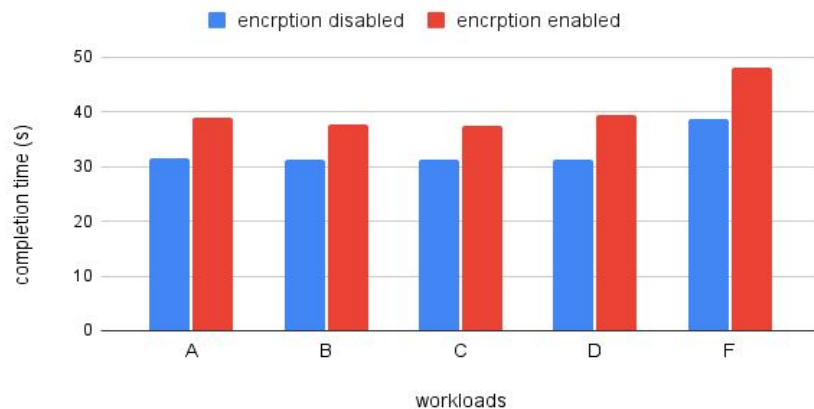
Evaluation: Value Encryption

Value encryption overhead in RocksDB



Lower is better


Value encryption overhead in Redis




Value encryption overhead is 24% for RocksDB and 23% for Redis.

Evaluation: Log Encryption

Log encryption overhead with RocksDB




Log encryption overhead with Redis




Log encryption of all workloads is 18% for RocksDB and 10% for Redis

Evaluation: Full GDPR Compliance Overhead

Full GDPR Compliance Overhead with RocksDB



Full GDPR Compliance Overhead with Redis



Lower is better

Full GDPR compliance overhead is 3.5x for RocksDB and 5x for Redis

Ways to reduce it: faster encryption algorithms and asynchronous logging

Current solutions to GDPR are **not** feasible

- Changes in application business logic
- Changes in database internals
- Varying performance overheads

GDPR Proxy:

- Generic GDPR compliant proxy for key-value stores
- No change in database internals
- Easily adaptable and extensible

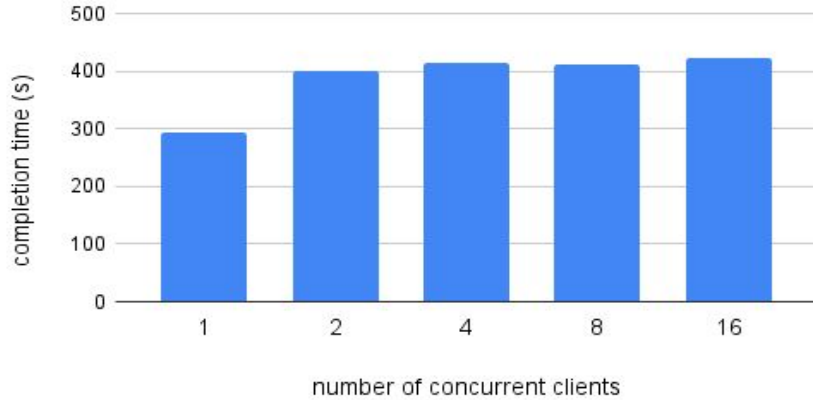


<https://github.com/ertugrulaypek/GDPRoxy>

Backup


Evaluation: Scaling

Scale Test with RocksDB



Lower is better

Scale Test with Redis



The completion time is not linearly increasing with concurrent clients

Evaluation: Disk usage

