

# UniBPF: Safe and Verifiable Unikernels Extensions

Kai-Chun Hsieh

Advisor: Masanori Misono

Chair of Computer Systems

<https://dse.in.tum.de/>



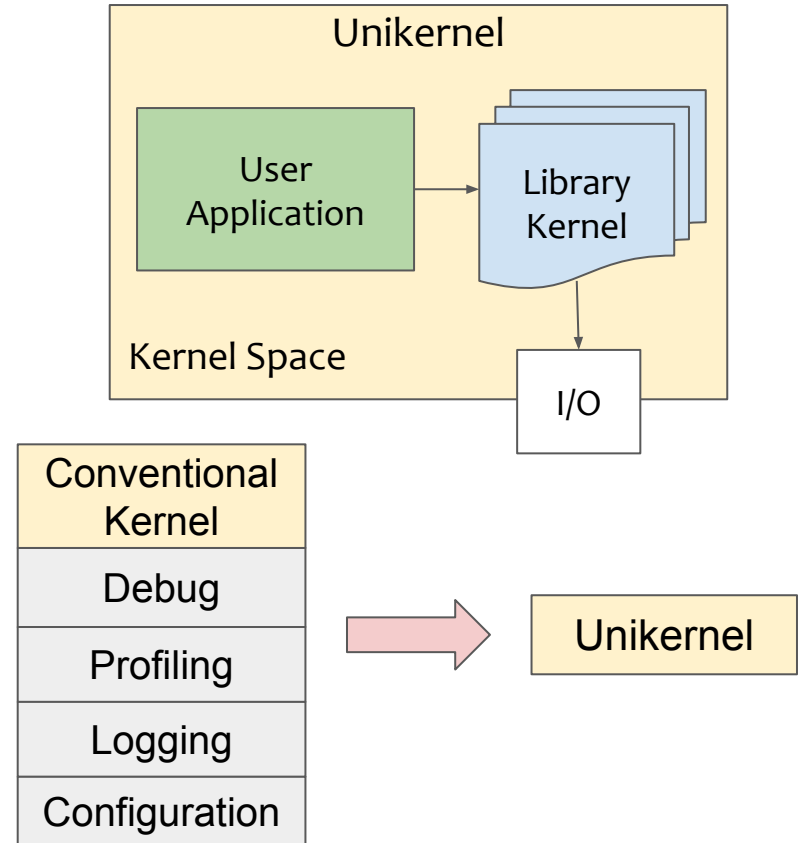
# Motivation

## Unikernels

- Kernel as a library
- Eliminate unneeded components.
- Optimize system procedures, e.g., system calls
- Compact, efficient, secure

## But...

- Lack of **debuggability**
- Lack of **observability**
- Lack of **runtime-extensibility**



## Extensible Unikernels with **BPF**:

- **eBPF Runtime** + kernel tracing with **interpreters**. But...



### Lack of verifier:

- Use an interpreter to provide sandboxed runtime



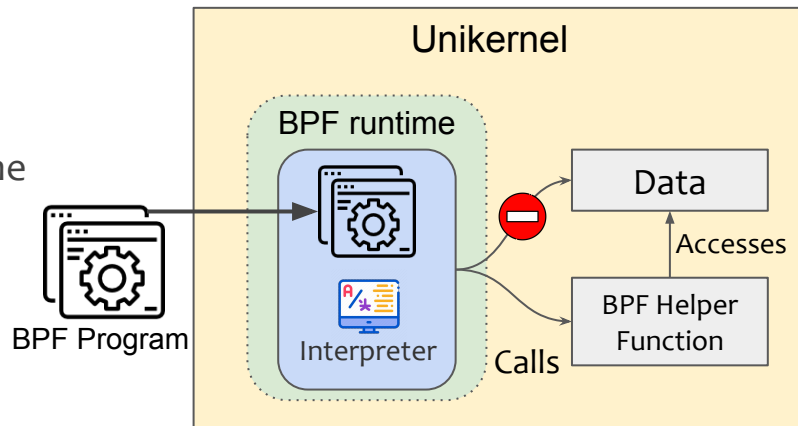
### Insufficient security guarantee:

- Cannot resist runtime errors



### Inefficient runtime:

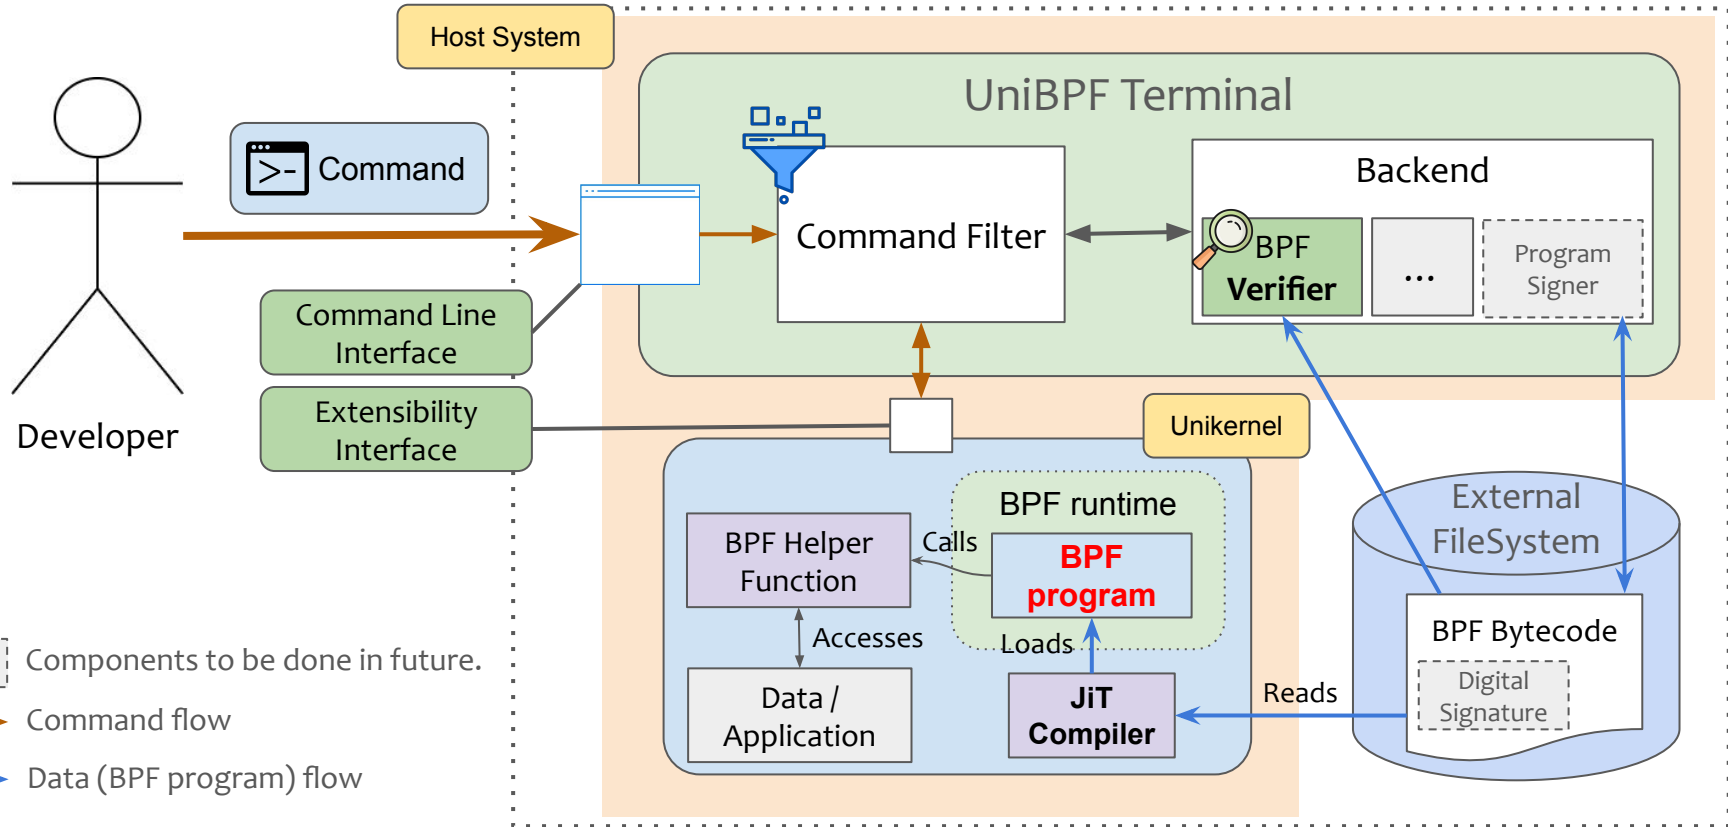
- Our work:  $\leq 600\%$  **slowdown** in instruction level v.s. JIT compiled



How can we have a safe and verifiable extension for Unikernels?

- Design Goal
  - **Safety:** Ensure safety of executing extension binaries
  - **Sustainable Design:** Easy to use, easy to maintain
  - **Performance:** Acceptable overhead and improve BPF runtime efficiency

# System Overview



# Background: extended Berkeley Packet Filter (eBPF)

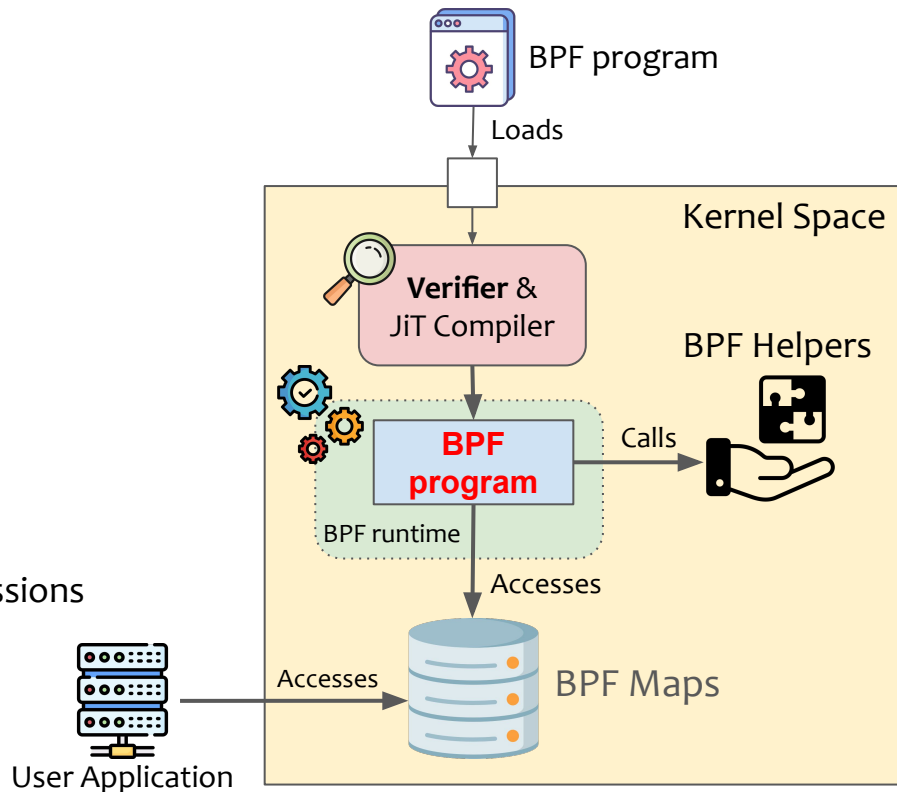
Lightweight in-kernel language VM

**Sandbox** property can be ensured by:

- Using **interpreters** (weaker)
- Using **verifiers** to verify in advance (stronger)
  - Detects potential sandbox escalation
  - Forbid undefined behaviors

Useful features:

- Maps (kv-store)
- Helper functions
- Program Types: Runtime context & helper permissions




# Outline



- ~~Motivation & Background~~
- Design Challenges
- Evaluation
- Further Ideas

1. Impact of Verification Processes on Unikernel Applications' Runtime
2. Feasibility of Integrating Verifier into Unikernel Application
3. Usability and Maintainability: Configuring Shared Verifier for Different Unikernels





# ① Verification can block Unikernel applications



Lack of **multi-processing** support:

- Application is the only process



Lack of comprehensive schedulers:

- CPU resource is released by voluntary “yields”



**Verification is time-consuming!**

- Our example BPF program: **12.05 ms** to verify 26 instructions.
- Lower-Bound: **8.82 ms**



With common approaches:

- Clients may experience huge **latencies**


```
1 __attribute__((section("executable"), used))
2 __u64 hash(uk_bpf_type_executable_t* context) {
3
4     __u64 sum = 0;
5
6     for(int index = 0; index < 256; index++) {
7         char* input = context->data + index;
8         if(input >= context->data_end) {
9             break;
10        }
11
12        char to_add = *input;
13
14        if(to_add >= 'A' && to_add <= 'Z') {
15            to_add += 'A' - 'a';
16        } else if(to_add >= '0' && to_add <= '9') {
17            to_add -= '0';
18        }
19
20        sum += to_add;
21    }
22
23    return sum;
24 }
```



Put BPF verifiers as processes on the host system where schedulers are more flexible

## ② BPF Verifiers Are Too Complicated to Integrate

- Common BPF verifiers are **complicated**:
  - PREVAIL (PLDI'19): 27,000 Lines of code
  - KLINT (NSDI'22): 13,000 Lines of code
- Common BPF verifiers need **complicated runtime**:
  - PREVAIL: C++ runtime library
  - KLINT: Python Interpreter
  - Linux BPF verifier: GPL License, Depends on Linux




Put BPF Verifiers on the host system utilizing the host system's runtime environment

### ③ Customizability Impedes Building a Unified Solution



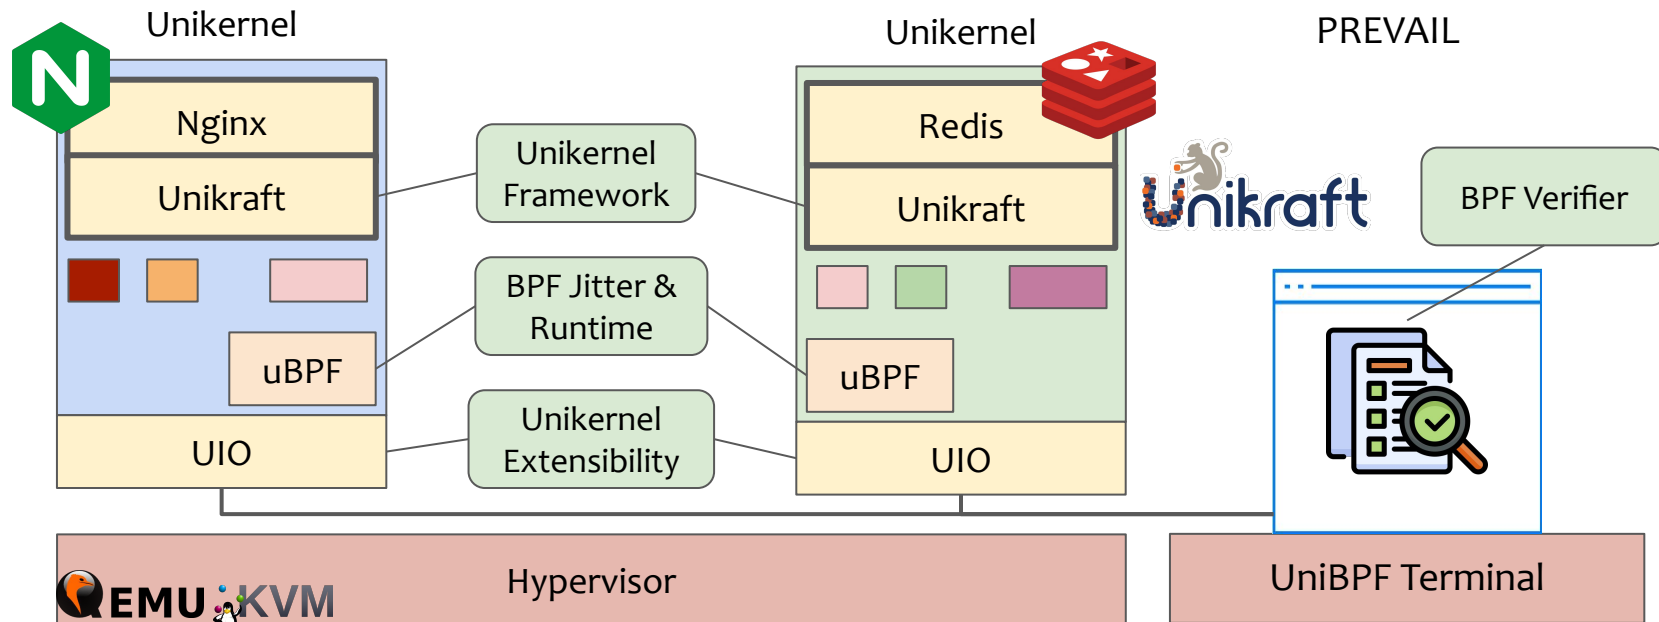
Our Goal: Maintain customizability for BPF runtime

- BPF Helper functions & program types
  - Keep compactness
  - Increase our system's usability
- But, without a standard framework:
    - Each Unikernel needs one BPF verifier: **Unmaintable!**
    - Waived support for customizable parts: Our work is **Meaningless!**



We provide libraries that allow developers to easily export their BPF runtime specifications

# Implementation




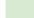
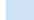

# Outline



- ~~Motivation & Background~~
- ~~Design Challenges~~
- Evaluation
- Further Ideas

# Evaluation - Safety

Evaluation Program	Result - Interpreter	Result - JiT Compiled	Result - UniBPF
OOB*	Terminated	Exploited	<b>Denied</b>
OOB* with Nullptr	Terminated	System crashed	<b>Denied</b>
Infinity Loop	System freezes	System freezes	<b>Denied</b>
Division by Zero	Error Ignored	Error Ignored	Partially <b>Denied</b>
Instruction Type Safety	Error Ignored	Error Ignored	<b>Denied</b>
Program Type Safety	Error Ignored	Error Ignored	<b>Denied</b>
Helper Function Type Safety	Error Ignored	Error Ignored	<b>Denied</b>

-  : Memory Safety
-  : Termination
-  : Runtime Errors
-  : Type Safety

UniBPF provide a safer BPF runtime extension for Unikernel

\*) OOB: Out of bound memory access

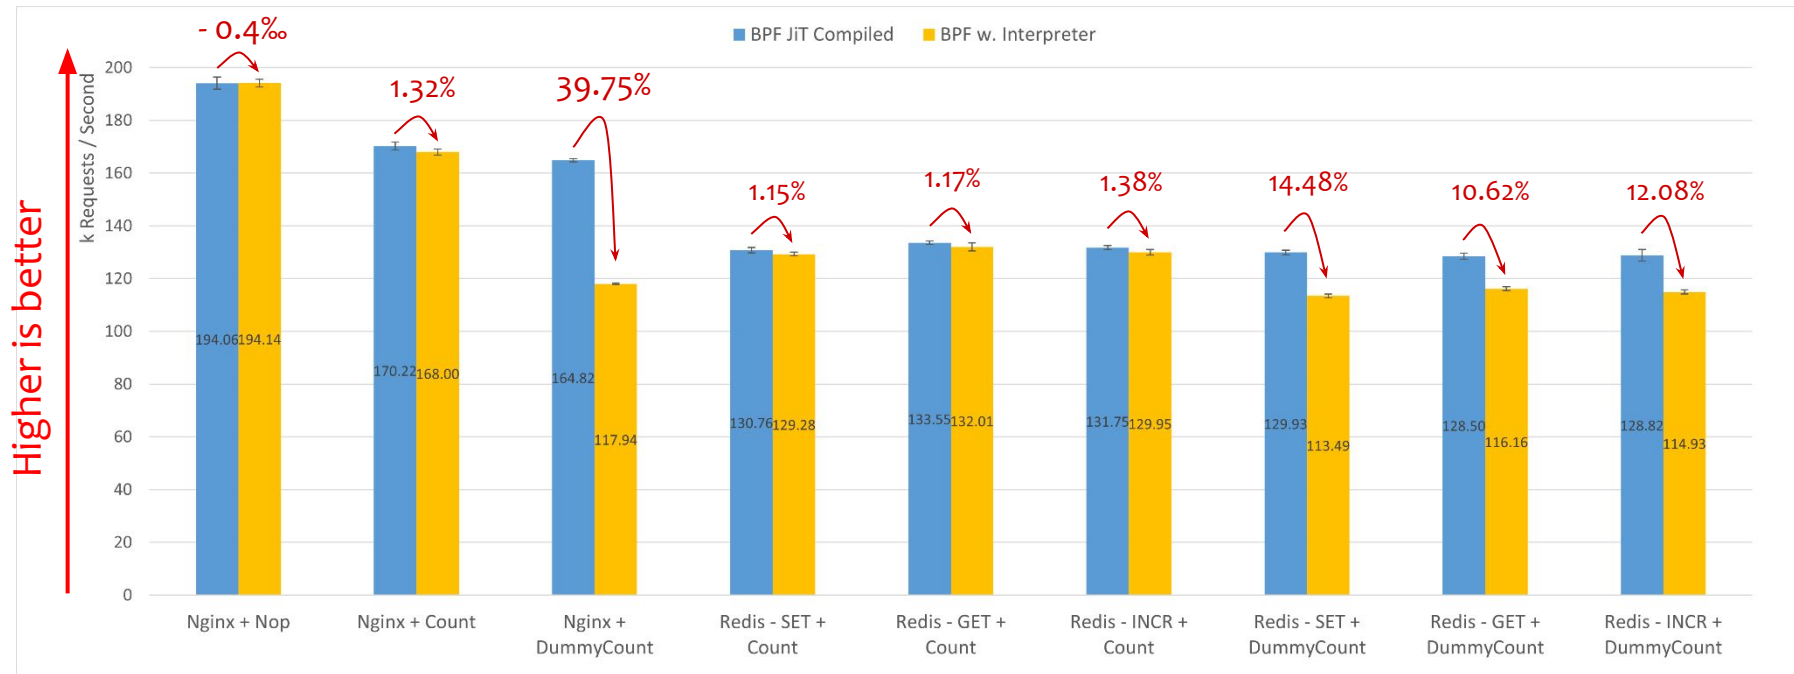
	Instructions	Verification Time Overhead*	Verification Memory Overhead*	JiT Time Overhead
Nop	2	8.82 ms	3328 kb	9.74 ms
Hash	26	12.05 ms (7.43 instr./ms)	4096 kb	9.79 ms
Adds	1002	43.60 ms (28.75 instr./ms)	5056 kb	9.85 ms

 : The lower bound overhead of the entire system.

\* : Overhead made to the host system.

The JiT compilation overhead and the corresponding verification overhead are negligible

# Evaluation - BPF Kernel Tracing Nginx and Redis




The improvement in jitted BPF runtime is more significant as the program size increases

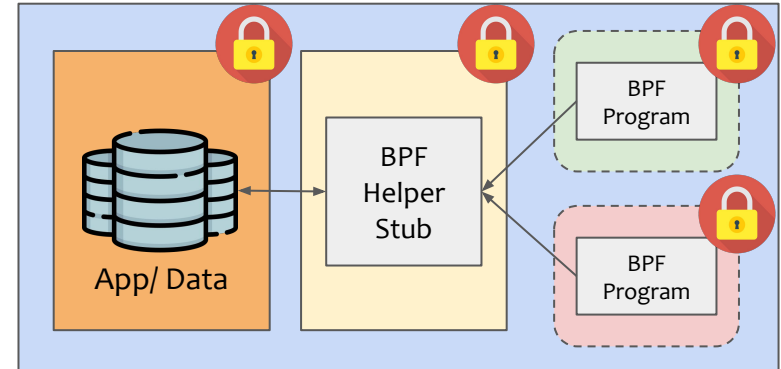


# Further Ideas

- Ensure verification integrity with **digital signature**
- More robust BPF runtime isolation:
  - Intel **MPK**
  - BPF helper function stub



- Support verification with BPF maps
- BPF program as configurations
- Secure verification process from malicious cloud provider: **Confidential VM**



- UniBPF provides **safer** BPF runtime
  - Resist runtime errors interpreters cannot
  - Protect jitted runtime from malicious codes
- Only brings **negligible overhead**
- Enables **more efficient runtime** through JiT compilation
  - Instruction level: Up to 600%
  - Kernel-Tracing:
    - Nginx: 40% ~
    - Redis: 14.48% ~

**Try it out!**

<https://github.com/TUM-DSE/ushell/>